

Convolutional Neural Networks (8DC00)

Navchetan Awasthi

Navchetan Awasthi Phd.



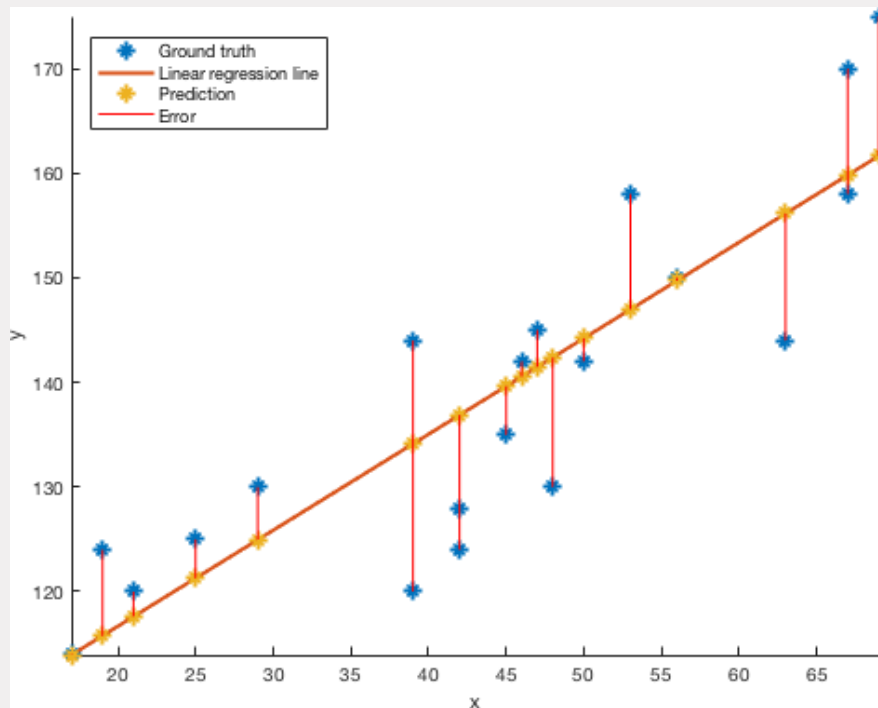
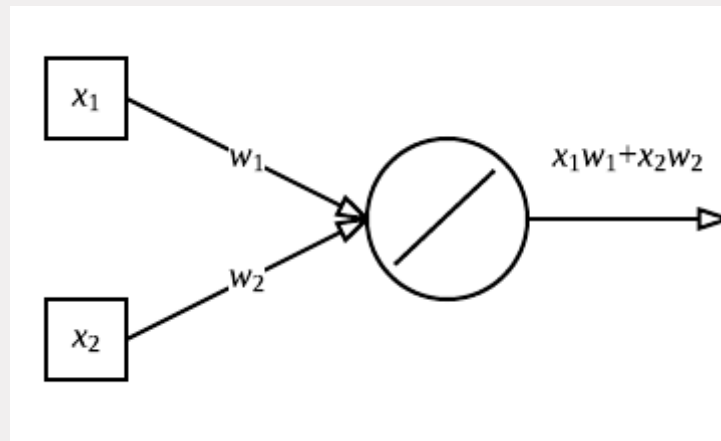
Background: B.Tech. Electronics and Communication Engineering, M.Tech. in Computational Science, PhD in Medical Imaging

Work/internship experience: Harvard Medical School, Massachusetts General Hospital, Indian Institute of Science

PhD Research: Deep learning, Medical image analysis, Medical Image reconstruction, Ultrasound Imaging, Photoacoustics, Inverse Problems

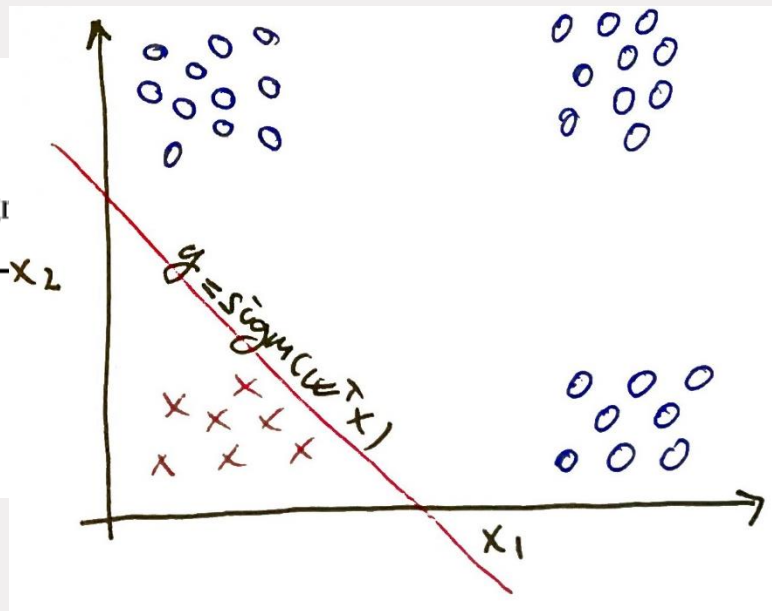
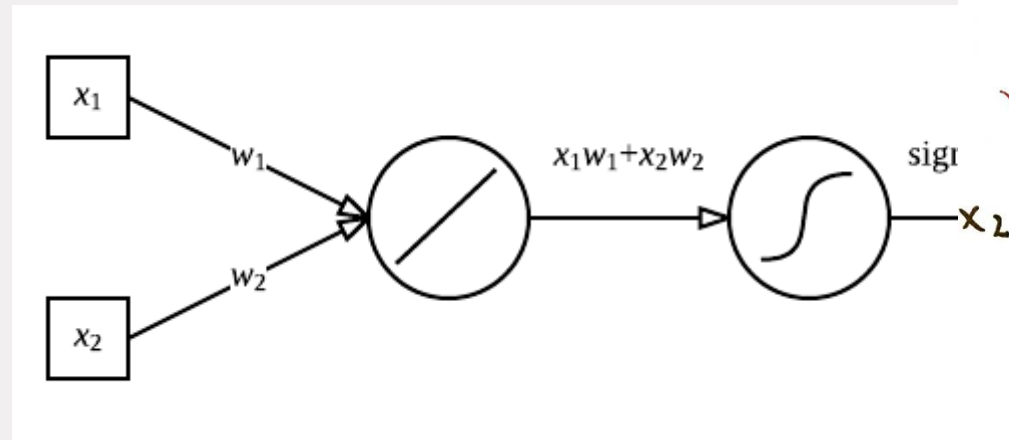
$$\hat{y} = \theta^T \mathbf{x}$$

Previously – Linear regression

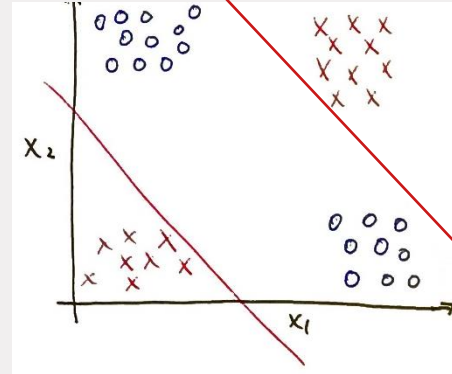
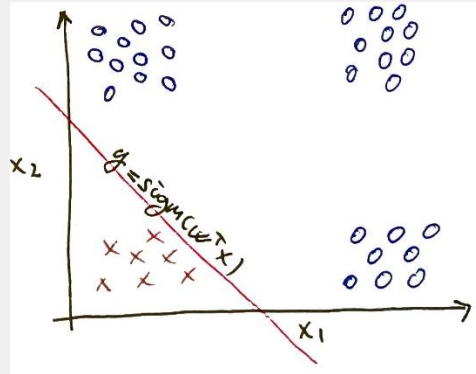
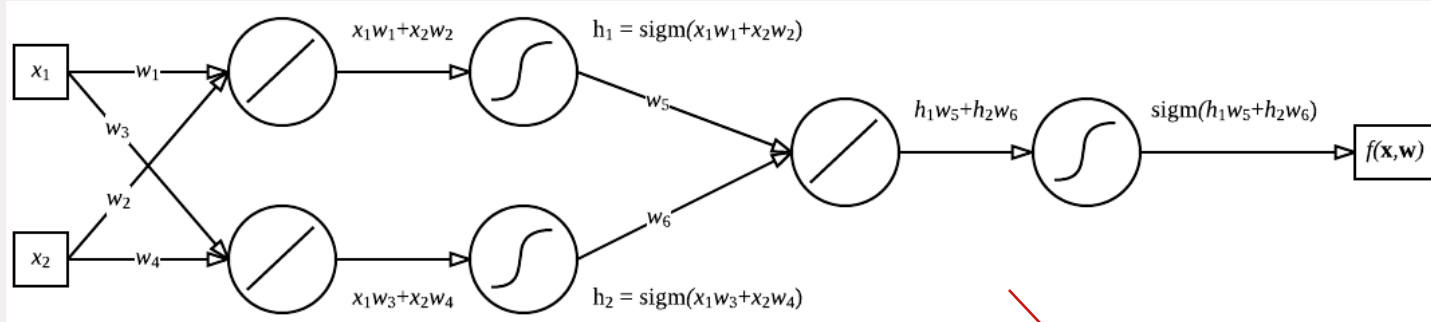


$$p(y = 1 \mid \mathbf{x}) = \sigma(\boldsymbol{\theta}^\top \mathbf{x})$$

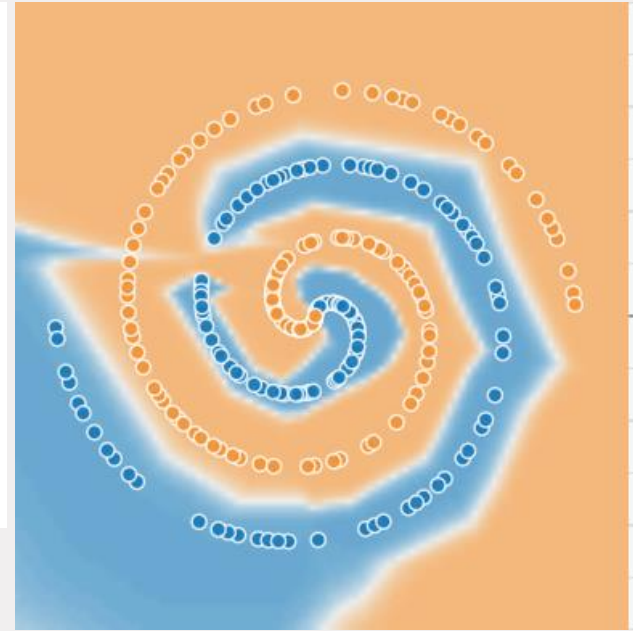
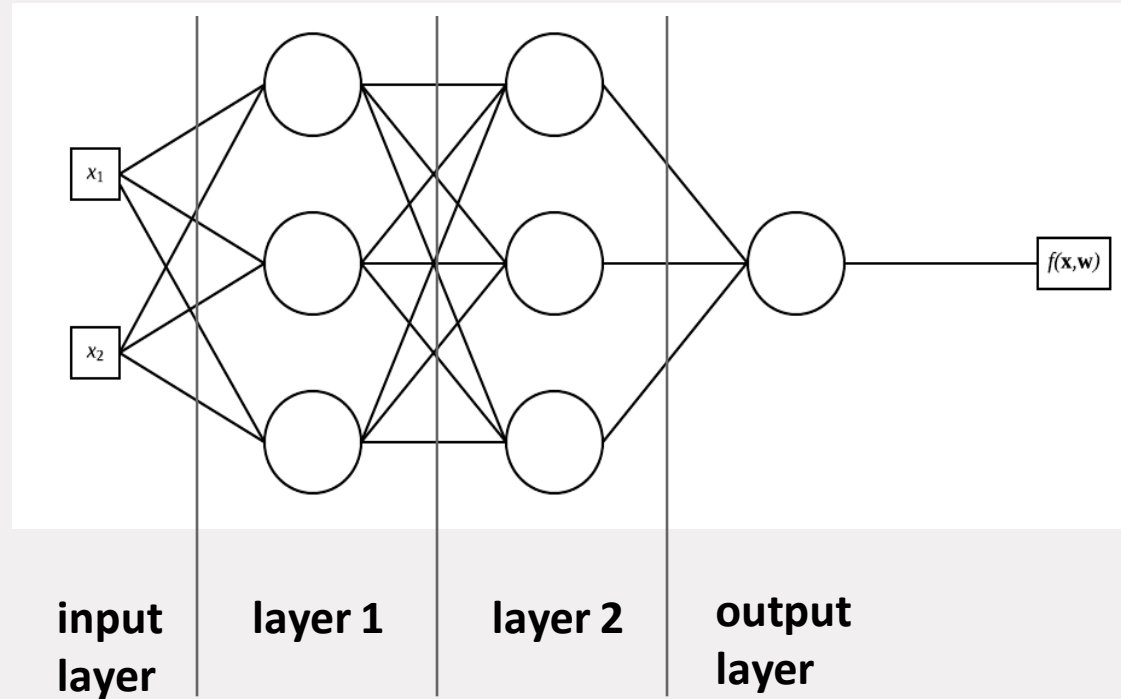
Previously – Logistic regression



Previously – Neural networks

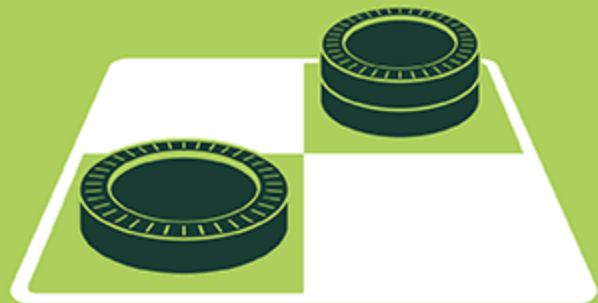


Previously – Neural networks



ARTIFICIAL INTELLIGENCE

Early artificial intelligence stirs excitement.



MACHINE LEARNING

Machine learning begins to flourish.



DEEP LEARNING

Deep learning breakthroughs drive AI boom.



Today: Convolutional neural networks (CNN)

- Neural networks → Convolutional neural networks

Building blocks for deep learning models for image analysis:

- Convolutional layer
- Max-pooling layer
- Not needed for the project, but will be on the exam

Learning outcomes

- Student can explain the concept of **convolutions** in a neural network
- Student can describe why we can use a **convolutional approach** for (medical) **images**
- Student can explain why convolutions enable development of **deep** (and large) **neural networks**
- Students can explain and apply the **max-pooling layer** in a convolutional neural network
- Students can motivate the choice for a **kernel size**

Lecture outline

- Images as input to neural network
- Reducing # of weights
- 1D convolutions
- 2D convolutions
- *Break (15 mins)*
- Kernels
- Max-pooling
- Interactive example

Images as input to neural networks



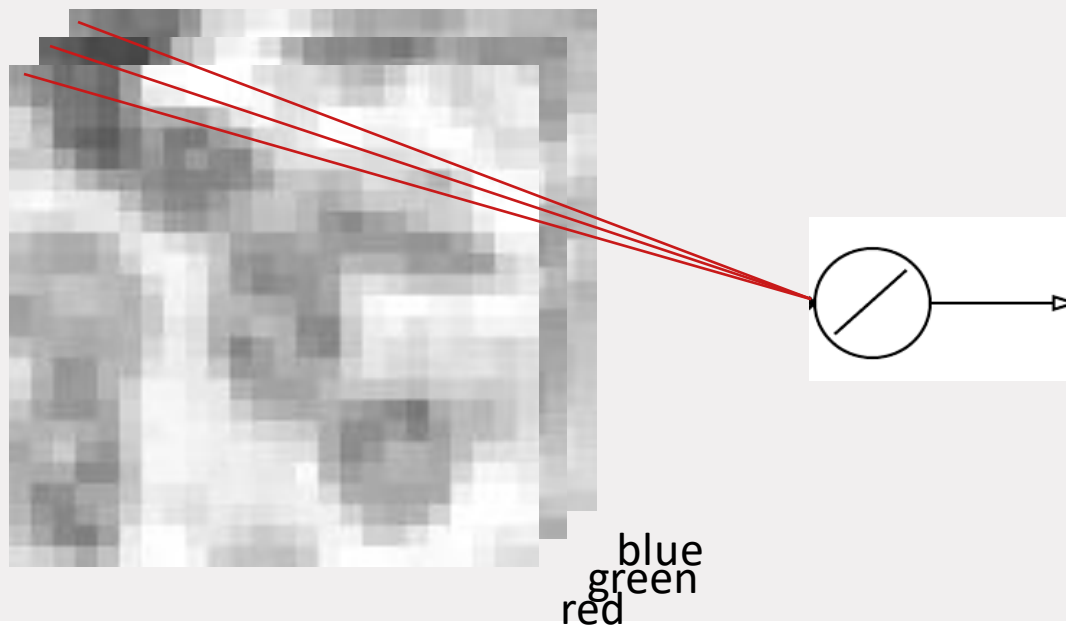
= 1728 features (\mathbf{x}_i are 1728 dimensional vectors)

If we train linear regression with these inputs (such as in the first practical), we will have 1728 weights w and a bias b .

$$\theta^T \mathbf{x}$$

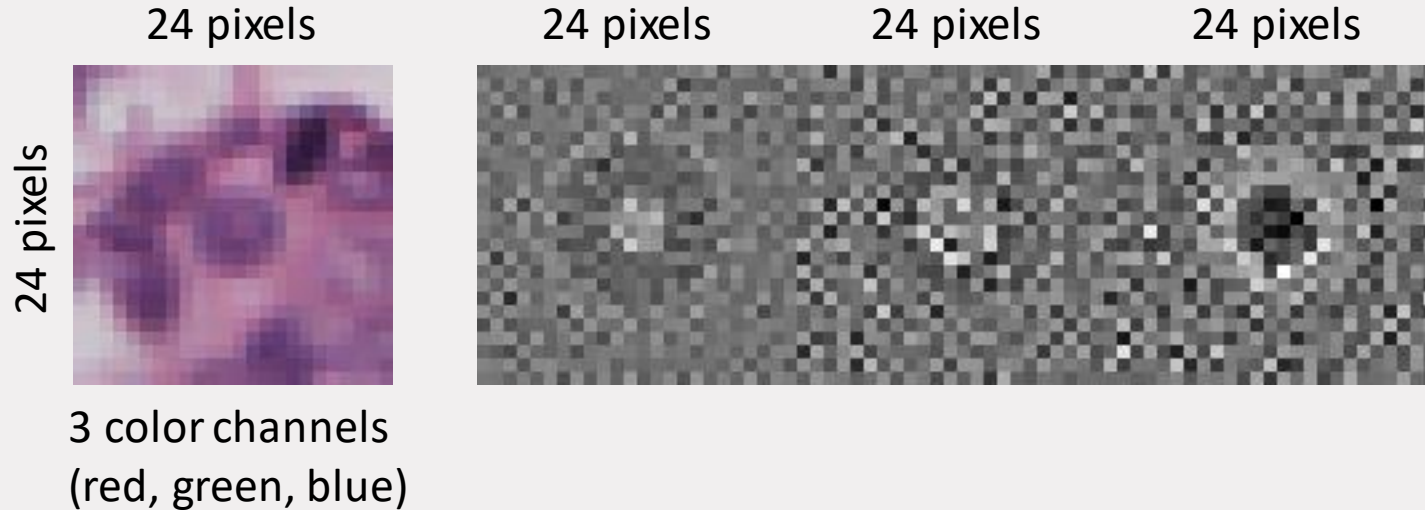
Every pixel is an input

Every pixel from every color channel is multiplied by a weight



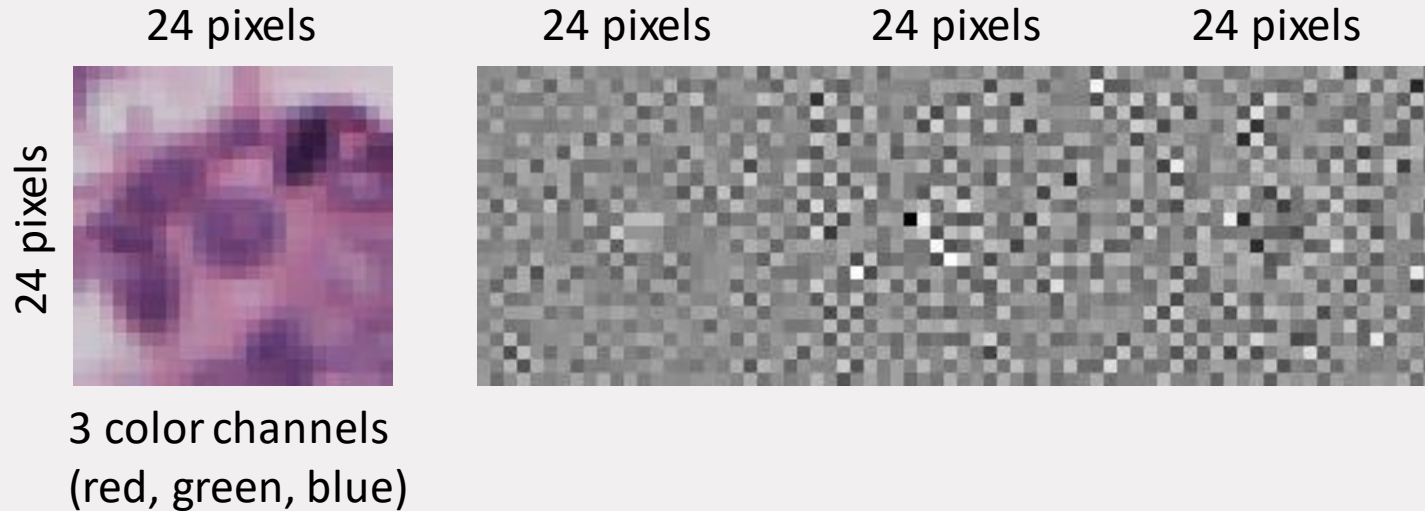
Post-training: visualizing what model has learned

Reshape vector of parameters into 24x24x3 image



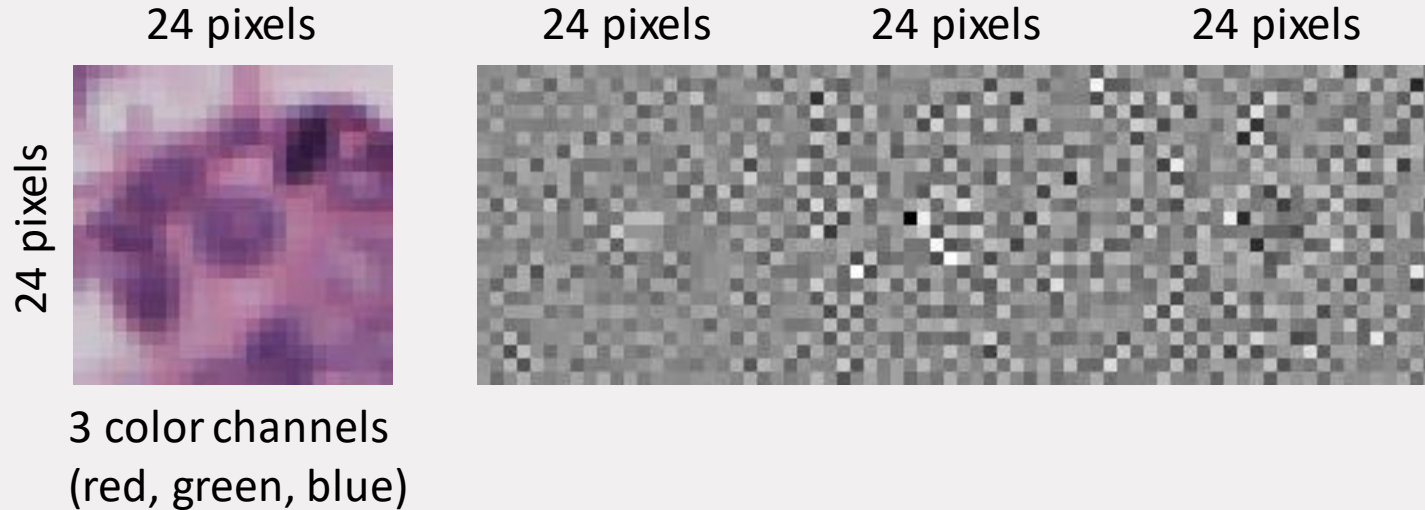
Post-training: visualizing what model has learned

Only 25% of training samples.. Looks noisy!

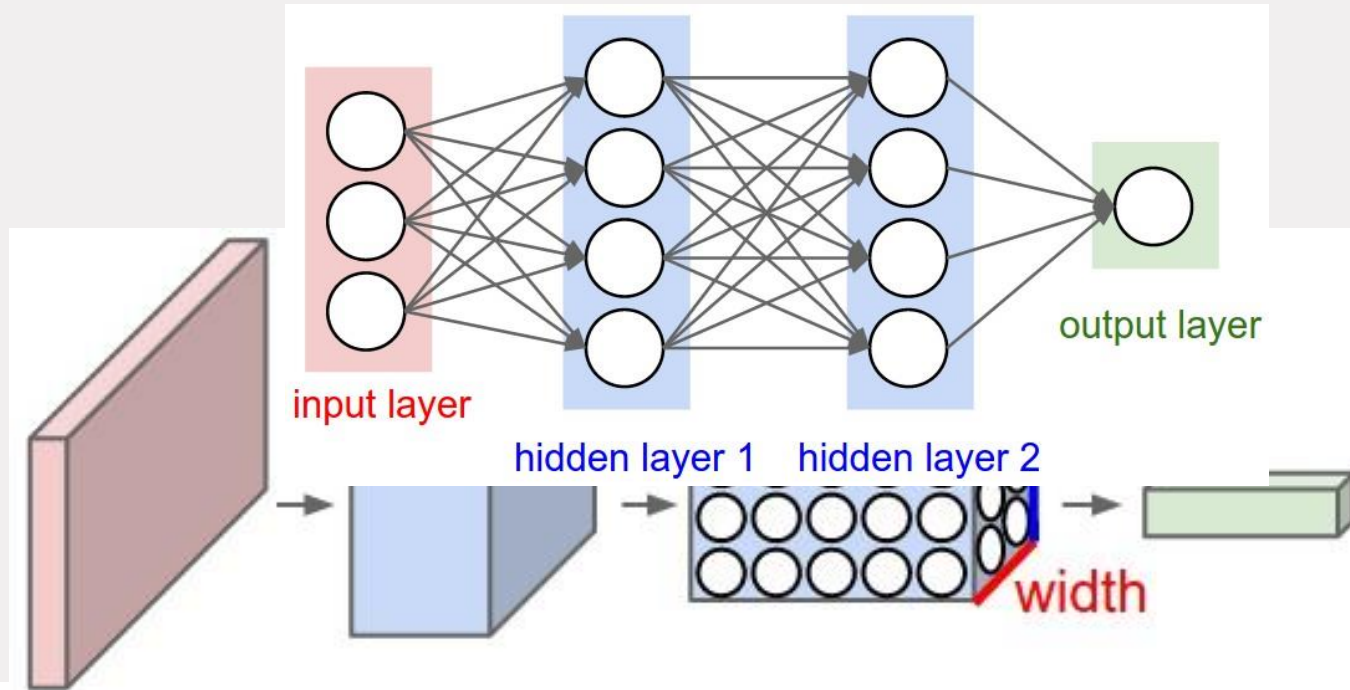


Post-training: visualizing what model has learned

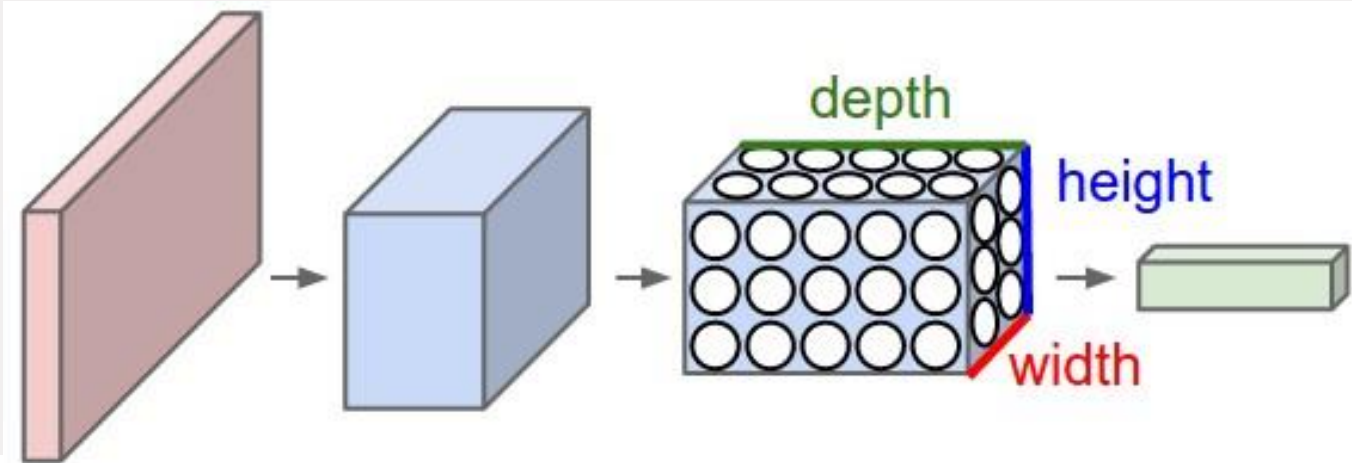
You can think of it as “there is not enough training data to reliably estimate all model weights”.



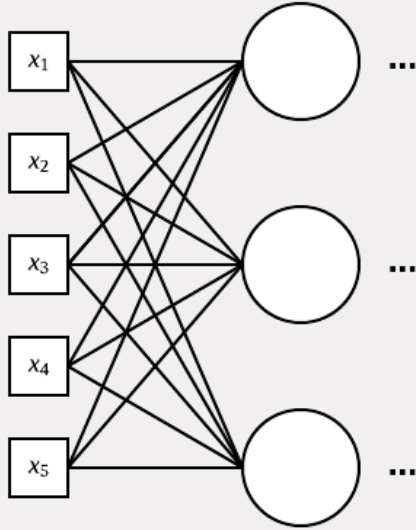
Many weights



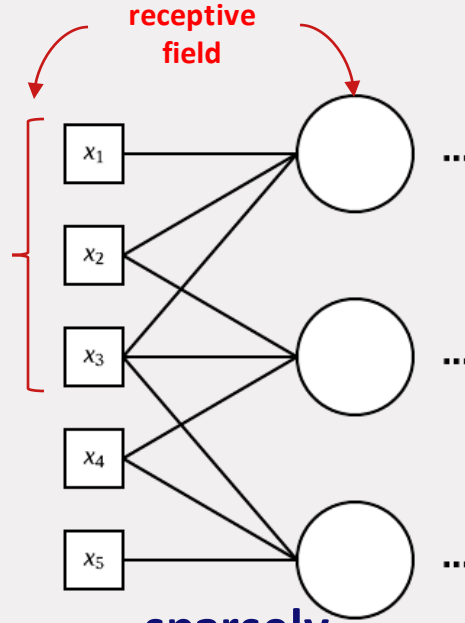
With large inputs such as images and deep networks, the number of weights “explodes”. We need a way to reduce the number of weights, without sacrificing performance.



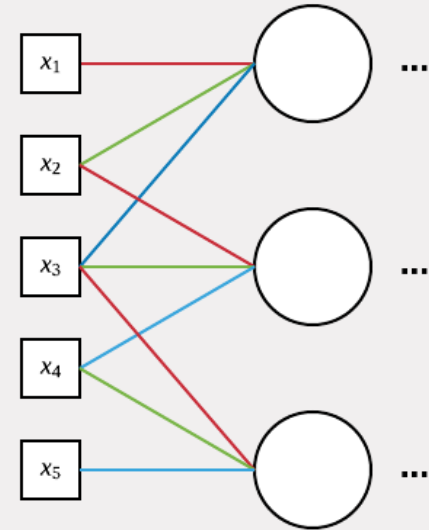
How many weights?



“regular” NN
15 weights

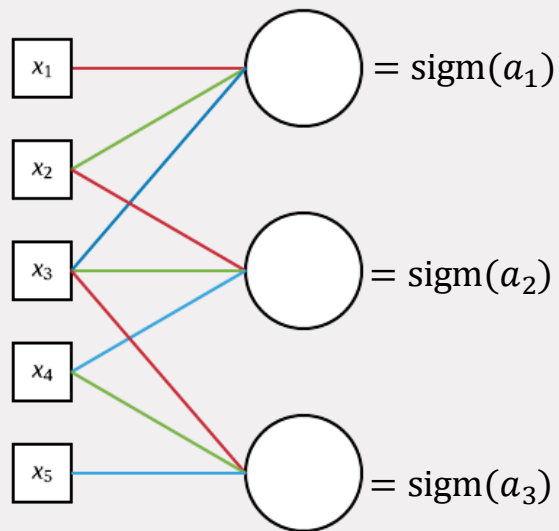


**sparsely
connected NN**
9 weights



shared weights
3 weights

Shared weights



shared weights
3 weights

$$a_1 = x_1 w_1 + x_2 w_2 + x_3 w_3$$

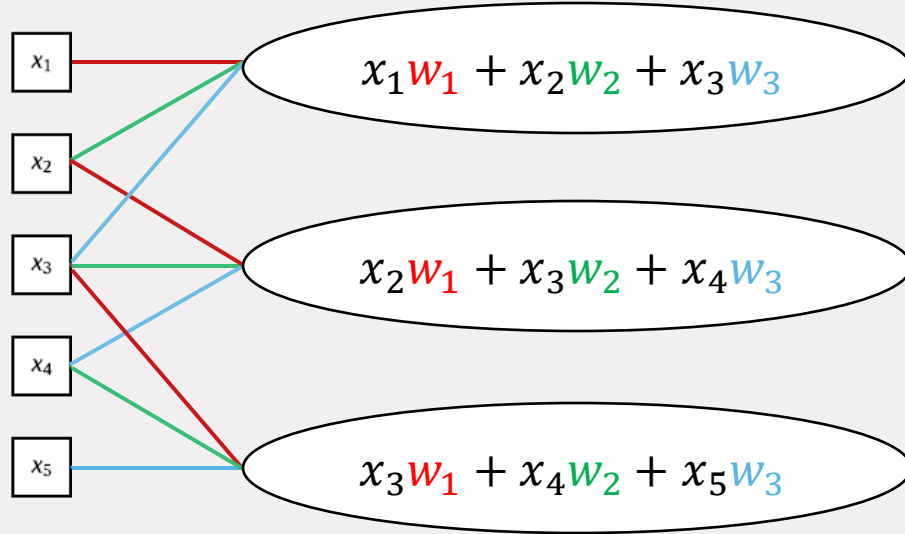
$$a_2 = x_2 w_1 + x_3 w_2 + x_4 w_3$$

$$a_3 = x_3 w_1 + x_4 w_2 + x_5 w_3$$

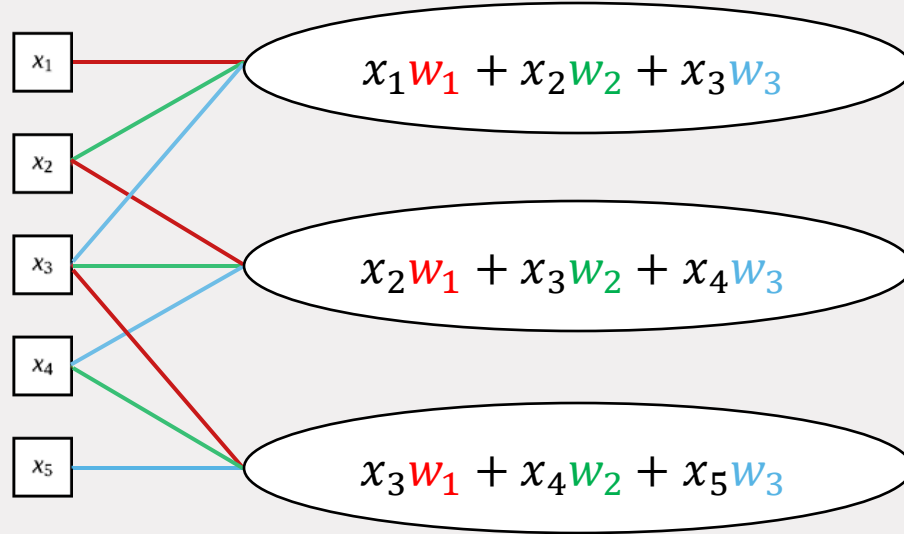
$$\begin{bmatrix} a_1 & a_2 & a_3 \end{bmatrix} = \begin{bmatrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{bmatrix} * \begin{bmatrix} w_1 & w_2 & w_3 \end{bmatrix}$$

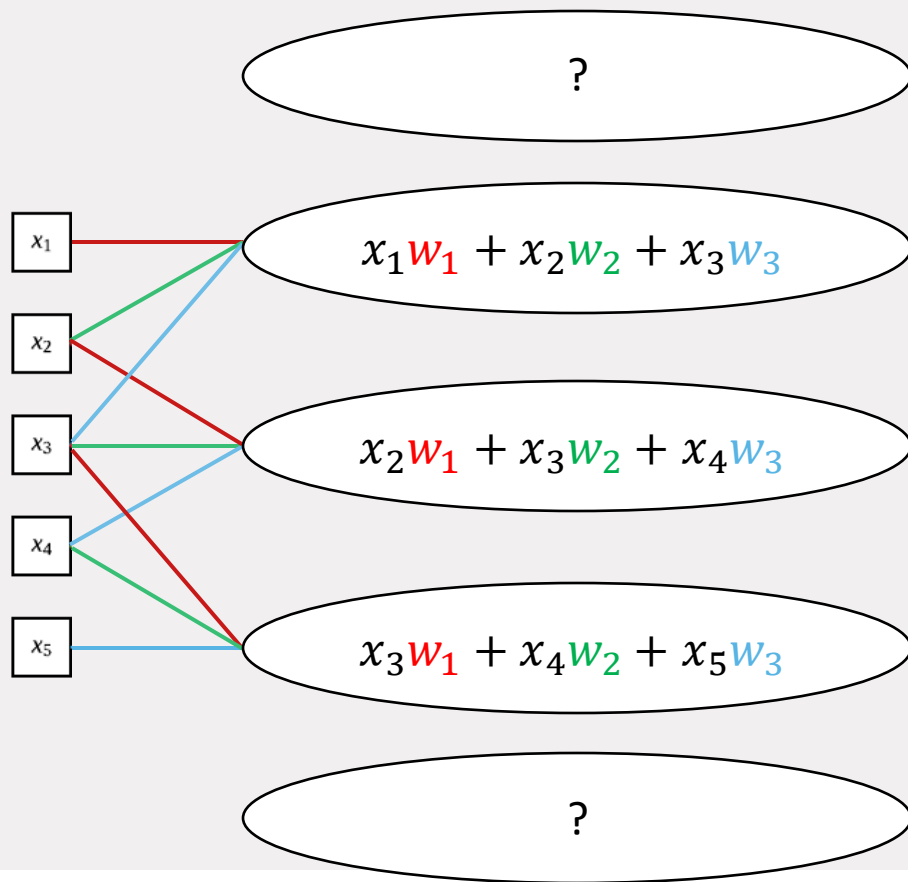
convolution, thus convolutional NN

1-D Convolution



1-D Convolution



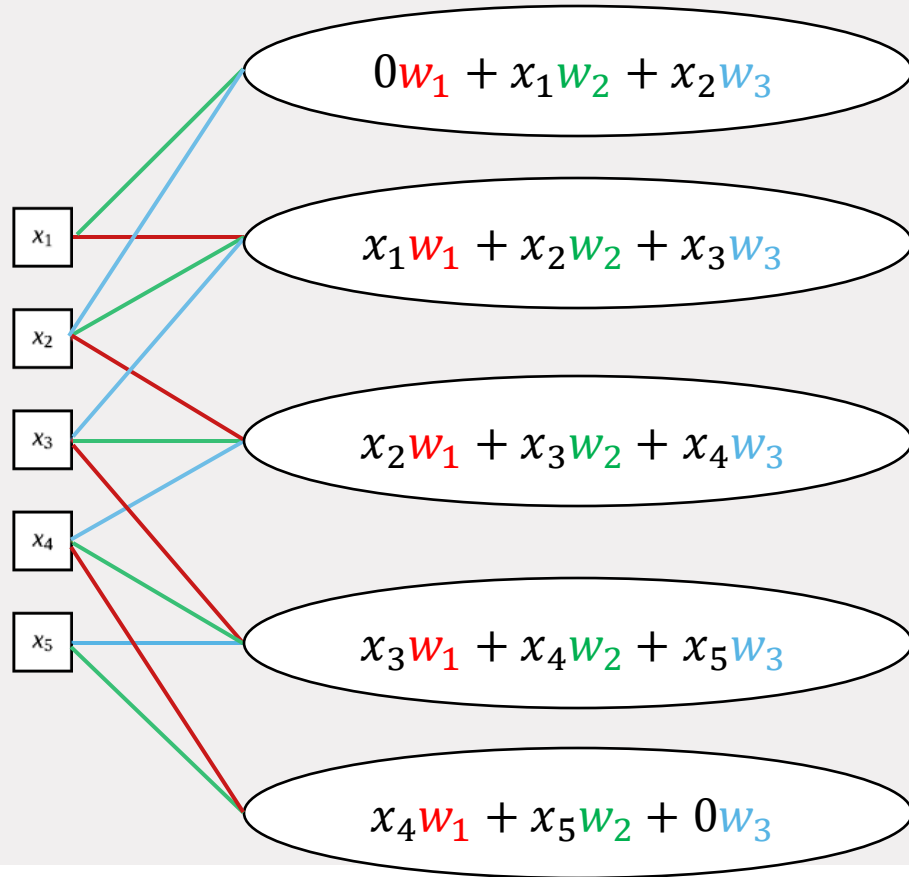


1-D Convolution

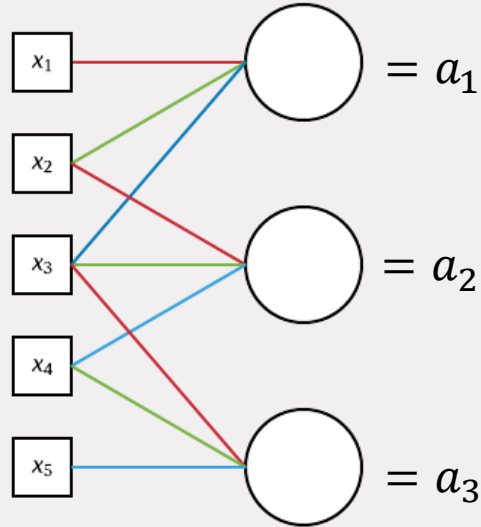
How can we keep the same number of features in hidden layer 1?

1-D Convolution

Zero-padding!



1-D Convolution



$$a_1 = x_1 w_1 + x_2 w_2 + x_3 w_3$$

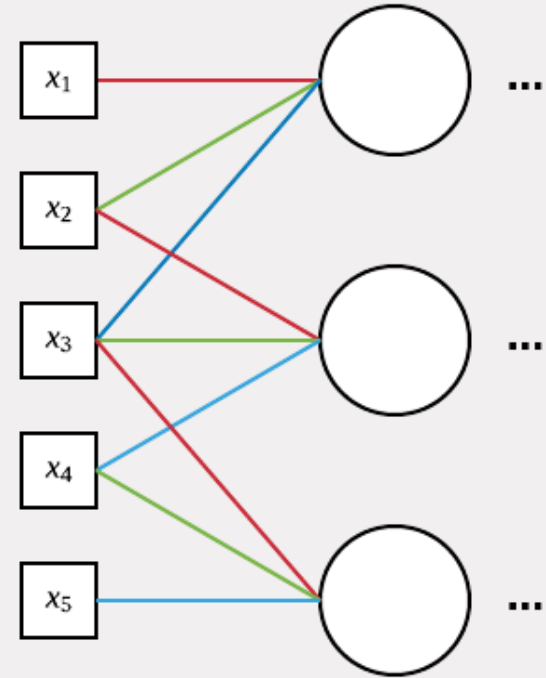
$$a_2 = x_2 w_1 + x_3 w_2 + x_4 w_3$$

$$a_3 = x_3 w_1 + x_4 w_2 + x_5 w_3$$

$$\begin{bmatrix} a_1 & a_2 & a_3 \end{bmatrix} = \begin{bmatrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{bmatrix} * \begin{bmatrix} w_1 & w_2 & w_3 \end{bmatrix}$$

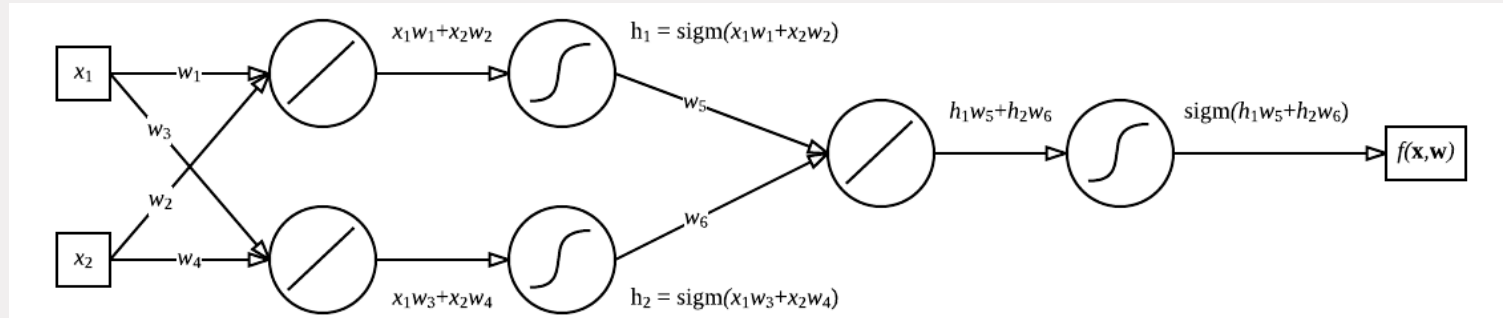
Properties of convolutional neural networks

- Sparse connectivity
- Weight sharing
- Parallel computations



Why does this work?

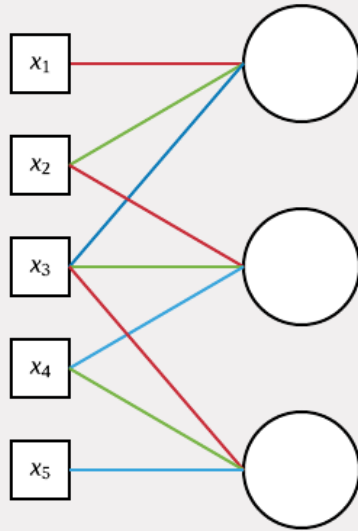
- Multiple layers
- Hidden layers contain features calculated from previous layers



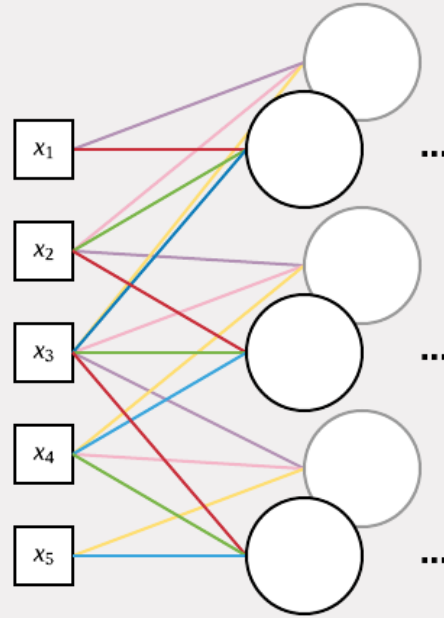
Why does this work?

- Different layers contain different transformation
 - Simple (e.g. edges, colors)
 - Complex (final layers)

One added benefit is that the learned transformations will be equivariant with translation (if the features/image is shifted up/down the features will still be detected).



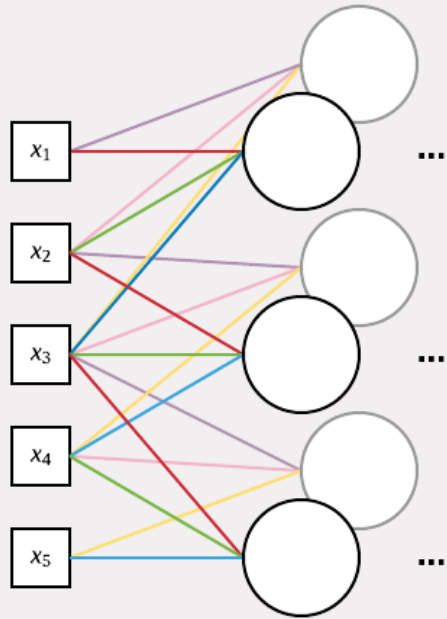
shared weights
3 weights



two sets shared weights
6 weights

We can add additional sets of weights that can learn additional interesting transformation of the input.

Note that the added neurons are not a new layer. They are part of layer 1.



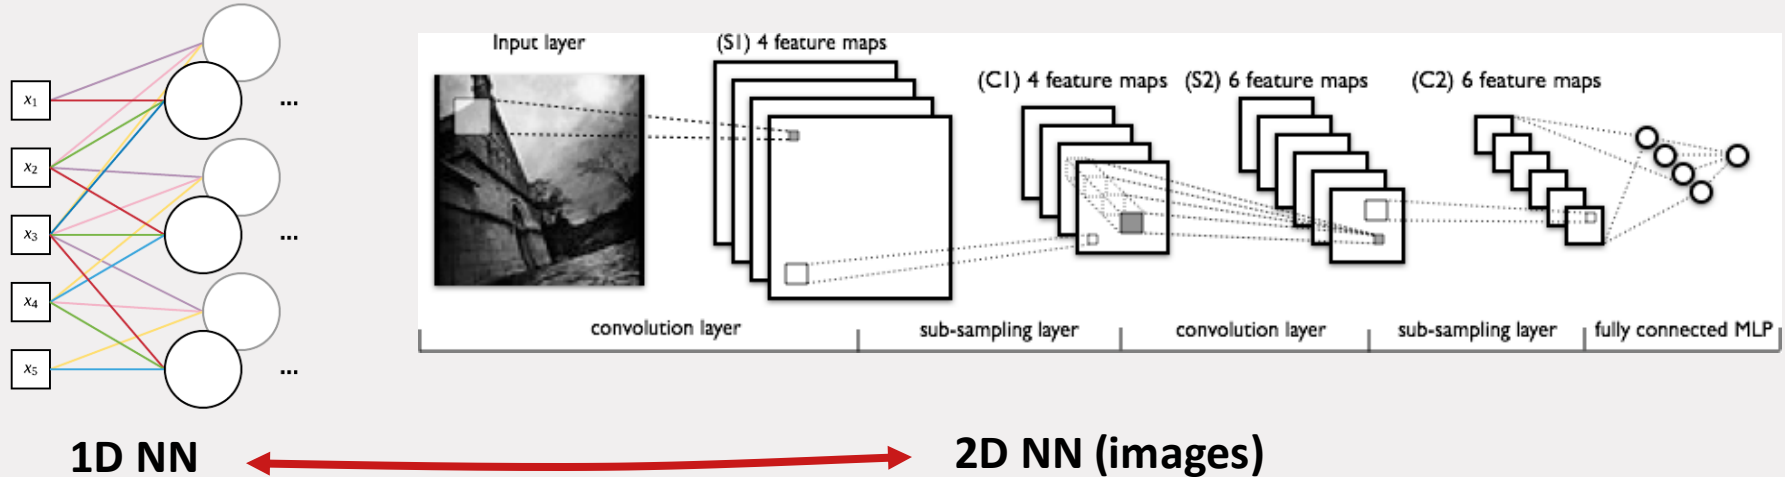
two sets shared weights
6 weights

$$\begin{bmatrix} a_{1,1} & a_{1,2} & a_{1,3} \end{bmatrix} = \begin{bmatrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{bmatrix} * \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \end{bmatrix}$$

$$\begin{bmatrix} a_{2,1} & a_{2,2} & a_{2,3} \end{bmatrix} = \begin{bmatrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{bmatrix} * \begin{bmatrix} w_{2,1} & w_{2,2} & w_{2,3} \end{bmatrix}$$

$\begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \end{bmatrix}$, and $\begin{bmatrix} w_{2,1} & w_{2,2} & w_{2,3} \end{bmatrix}$ are **convolution kernels**. They extract features. However, they are not hand-designed features – they were learned by the neural network.

Convolutional neural networks are ideal for images



Because of the weight sharing, convolutional neural networks only work with structured data (such as images) as inputs.

$$\begin{array}{|c|c|c|c|c|} \hline I_1 & I_2 & I_3 & I_4 & I_5 \\ \hline I_6 & I_7 & I_8 & I_9 & I_{10} \\ \hline I_{11} & I_{12} & I_{13} & I_{14} & I_{15} \\ \hline I_{16} & I_{17} & I_{18} & I_{19} & I_{20} \\ \hline I_{21} & I_{22} & I_{23} & I_{24} & I_{25} \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline w_1 & w_2 & w_3 \\ \hline w_4 & w_5 & w_6 \\ \hline w_7 & w_8 & w_9 \\ \hline \end{array} =$$

$$\begin{array}{|c|c|c|} \hline I_1 w_1 + I_2 w_2 + I_3 w_3 & & \\ \hline + & & \\ I_6 w_4 + I_7 w_5 + I_8 w_6 & & \\ \hline + & & \\ I_{11} w_7 + I_{12} w_8 + I_{13} w_9 & & \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|c|c|} \hline I_1 & I_2 & I_3 & I_4 & I_5 \\ \hline I_6 & I_7 & I_8 & I_9 & I_{10} \\ \hline I_{11} & I_{12} & I_{13} & I_{14} & I_{15} \\ \hline I_{16} & I_{17} & I_{18} & I_{19} & I_{20} \\ \hline I_{21} & I_{22} & I_{23} & I_{24} & I_{25} \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline w_1 & w_2 & w_3 \\ \hline w_4 & w_5 & w_6 \\ \hline w_7 & w_8 & w_9 \\ \hline \end{array} =$$

$$\begin{array}{|c|c|c|} \hline \begin{array}{c} I_1 w_1 + I_2 w_2 + I_3 w_3 \\ + \\ I_6 w_4 + I_7 w_5 + I_8 w_6 \\ + \\ I_{11} w_7 + I_{12} w_8 + I_{13} w_9 \end{array} & \begin{array}{c} I_2 w_1 + I_3 w_2 + I_4 w_3 \\ + \\ I_7 w_4 + I_8 w_5 + I_9 w_6 \\ + \\ I_{12} w_7 + I_{13} w_8 + I_{14} w_9 \end{array} & \\ \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|c|c|} \hline I_1 & I_2 & I_3 & I_4 & I_5 \\ \hline I_6 & I_7 & I_8 & I_9 & I_{10} \\ \hline I_{11} & I_{12} & I_{13} & I_{14} & I_{15} \\ \hline I_{16} & I_{17} & I_{18} & I_{19} & I_{20} \\ \hline I_{21} & I_{22} & I_{23} & I_{24} & I_{25} \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline w_1 & w_2 & w_3 \\ \hline w_4 & w_5 & w_6 \\ \hline w_7 & w_8 & w_9 \\ \hline \end{array} =$$

$$\begin{array}{|c|c|c|} \hline I_1 w_1 + I_2 w_2 + I_3 w_3 & I_2 w_1 + I_3 w_2 + I_4 w_3 & I_3 w_1 + I_4 w_2 + I_5 w_3 \\ \hline + & + & + \\ I_6 w_4 + I_7 w_5 + I_8 w_6 & I_7 w_4 + I_8 w_5 + I_9 w_6 & I_8 w_4 + I_9 w_5 + I_{10} w_6 \\ \hline + & + & + \\ I_{11} w_7 + I_{12} w_8 + I_{13} w_9 & I_{12} w_7 + I_{13} w_8 + I_{14} w_9 & I_{13} w_7 + I_{14} w_8 + I_{15} w_9 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|c|c|} \hline I_1 & I_2 & I_3 & I_4 & I_5 \\ \hline I_6 & I_7 & I_8 & I_9 & I_{10} \\ \hline I_{11} & I_{12} & I_{13} & I_{14} & I_{15} \\ \hline I_{16} & I_{17} & I_{18} & I_{19} & I_{20} \\ \hline I_{21} & I_{22} & I_{23} & I_{24} & I_{25} \\ \hline \end{array}
 *
 \begin{array}{|c|c|c|} \hline w_1 & w_2 & w_3 \\ \hline w_4 & w_5 & w_6 \\ \hline w_7 & w_8 & w_9 \\ \hline \end{array}
 =$$

$$\begin{array}{|c|c|c|} \hline I_1 w_1 + I_2 w_2 + I_3 w_3 + I_6 w_4 + I_7 w_5 + I_8 w_6 + I_{11} w_7 + I_{12} w_8 + I_{13} w_9 \\ \hline I_2 w_1 + I_3 w_2 + I_4 w_3 + I_7 w_4 + I_8 w_5 + I_9 w_6 + I_{12} w_7 + I_{13} w_8 + I_{14} w_9 \\ \hline I_3 w_1 + I_4 w_2 + I_5 w_3 + I_8 w_4 + I_9 w_5 + I_{10} w_6 + I_{13} w_7 + I_{14} w_8 + I_{15} w_9 \\ \hline I_6 w_1 + I_7 w_2 + I_8 w_3 + I_{11} w_4 + I_{12} w_5 + I_{13} w_6 + I_{16} w_7 + I_{17} w_8 + I_{18} w_9 \\ \hline \\ \hline \\ \hline \end{array}$$

I_1	I_2	I_3	I_4	I_5
I_6	I_7	I_8	I_9	I_{10}
I_{11}	I_{12}	I_{13}	I_{14}	I_{15}
I_{16}	I_{17}	I_{18}	I_{19}	I_{20}
I_{21}	I_{22}	I_{23}	I_{24}	I_{25}

*

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

=

$I_1 w_1 + I_2 w_2 + I_3 w_3$ $+$ $I_6 w_4 + I_7 w_5 + I_8 w_6$ $+$ $I_{11} w_7 + I_{12} w_8 + I_{13} w_9$	$I_2 w_1 + I_3 w_2 + I_4 w_3$ $+$ $I_7 w_4 + I_8 w_5 + I_9 w_6$ $+$ $I_{12} w_7 + I_{13} w_8 + I_{14} w_9$	$I_3 w_1 + I_4 w_2 + I_5 w_3$ $+$ $I_8 w_4 + I_9 w_5 + I_{10} w_6$ $+$ $I_{13} w_7 + I_{14} w_8 + I_{15} w_9$
$I_6 w_1 + I_7 w_2 + I_8 w_3$ $+$ $I_{11} w_4 + I_{12} w_5 + I_{13} w_6$ $+$ $I_{16} w_7 + I_{17} w_8 + I_{18} w_9$	$I_7 w_1 + I_8 w_2 + I_9 w_3$ $+$ $I_{12} w_4 + I_{13} w_5 + I_{14} w_6$ $+$ $I_{17} w_7 + I_{18} w_8 + I_{19} w_9$	

$$\begin{array}{|c|c|c|c|c|} \hline I_1 & I_2 & I_3 & I_4 & I_5 \\ \hline I_6 & I_7 & I_8 & I_9 & I_{10} \\ \hline I_{11} & I_{12} & I_{13} & I_{14} & I_{15} \\ \hline I_{16} & I_{17} & I_{18} & I_{19} & I_{20} \\ \hline I_{21} & I_{22} & I_{23} & I_{24} & I_{25} \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline w_1 & w_2 & w_3 \\ \hline w_4 & w_5 & w_6 \\ \hline w_7 & w_8 & w_9 \\ \hline \end{array} =$$

$$\begin{array}{|c|c|c|} \hline I_1 w_1 + I_2 w_2 + I_3 w_3 + I_6 w_4 + I_7 w_5 + I_8 w_6 + I_{11} w_7 + I_{12} w_8 + I_{13} w_9 \\ \hline I_2 w_1 + I_3 w_2 + I_4 w_3 + I_7 w_4 + I_8 w_5 + I_9 w_6 + I_{12} w_7 + I_{13} w_8 + I_{14} w_9 \\ \hline I_3 w_1 + I_4 w_2 + I_5 w_3 + I_8 w_4 + I_9 w_5 + I_{10} w_6 + I_{13} w_7 + I_{14} w_8 + I_{15} w_9 \\ \hline I_6 w_1 + I_7 w_2 + I_8 w_3 + I_{11} w_4 + I_{12} w_5 + I_{13} w_6 + I_{16} w_7 + I_{17} w_8 + I_{18} w_9 \\ \hline I_7 w_1 + I_8 w_2 + I_9 w_3 + I_{12} w_4 + I_{13} w_5 + I_{14} w_6 + I_{17} w_7 + I_{18} w_8 + I_{19} w_9 \\ \hline I_8 w_1 + I_9 w_2 + I_{10} w_3 + I_{13} w_4 + I_{14} w_5 + I_{15} w_6 + I_{18} w_7 + I_{19} w_8 + I_{20} w_9 \\ \hline \end{array}$$

$$\begin{array}{ccccc}
 I_1 & I_2 & I_3 & I_4 & I_5 \\
 I_6 & I_7 & I_8 & I_9 & I_{10} \\
 I_{11} & I_{12} & I_{13} & I_{14} & I_{15} \\
 I_{16} & I_{17} & I_{18} & I_{19} & I_{20} \\
 I_{21} & I_{22} & I_{23} & I_{24} & I_{25}
 \end{array}
 *
 \begin{array}{ccc}
 w_1 & w_2 & w_3 \\
 w_4 & w_5 & w_6 \\
 w_7 & w_8 & w_9
 \end{array}
 =$$

$$\begin{array}{ccc}
 \begin{array}{c} I_1 w_1 + I_2 w_2 + I_3 w_3 \\ + \\ I_6 w_4 + I_7 w_5 + I_8 w_6 \\ + \\ I_{11} w_7 + I_{12} w_8 + I_{13} w_9 \end{array} &
 \begin{array}{c} I_2 w_1 + I_3 w_2 + I_4 w_3 \\ + \\ I_7 w_4 + I_8 w_5 + I_9 w_6 \\ + \\ I_{12} w_7 + I_{13} w_8 + I_{14} w_9 \end{array} &
 \begin{array}{c} I_3 w_1 + I_4 w_2 + I_5 w_3 \\ + \\ I_8 w_4 + I_9 w_5 + I_{10} w_6 \\ + \\ I_{13} w_7 + I_{14} w_8 + I_{15} w_9 \end{array} \\
 \begin{array}{c} I_6 w_1 + I_7 w_2 + I_8 w_3 \\ + \\ I_{11} w_4 + I_{12} w_5 + I_{13} w_6 \\ + \\ I_{16} w_7 + I_{17} w_8 + I_{18} w_9 \end{array} &
 \begin{array}{c} I_7 w_1 + I_8 w_2 + I_9 w_3 \\ + \\ I_{12} w_4 + I_{13} w_5 + I_{14} w_6 \\ + \\ I_{17} w_7 + I_{18} w_8 + I_{19} w_9 \end{array} &
 \begin{array}{c} I_8 w_1 + I_9 w_2 + I_{10} w_3 \\ + \\ I_{13} w_4 + I_{14} w_5 + I_{15} w_6 \\ + \\ I_{18} w_7 + I_{19} w_8 + I_{20} w_9 \end{array} \\
 \begin{array}{c} I_{11} w_1 + I_{12} w_2 + I_{13} w_3 \\ + \\ I_{16} w_4 + I_{17} w_5 + I_{18} w_6 \\ + \\ I_{21} w_7 + I_{22} w_8 + I_{23} w_9 \end{array} &
 \begin{array}{c} \\ \\ \\ \end{array} &
 \begin{array}{c} \\ \\ \\ \end{array}
 \end{array}$$

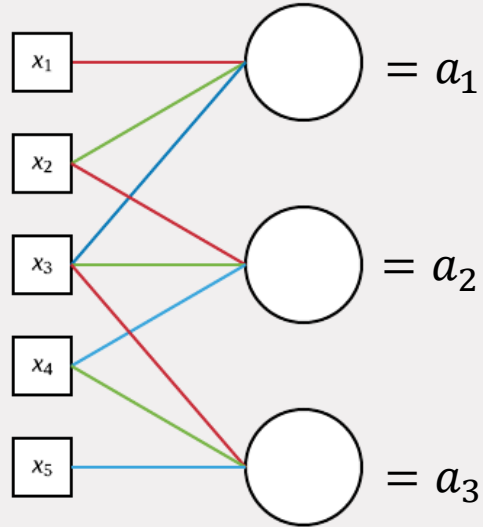
$$\begin{array}{ccccc}
 I_1 & I_2 & I_3 & I_4 & I_5 \\
 I_6 & I_7 & I_8 & I_9 & I_{10} \\
 I_{11} & I_{12} & I_{13} & I_{14} & I_{15} \\
 I_{16} & I_{17} & I_{18} & I_{19} & I_{20} \\
 I_{21} & I_{22} & I_{23} & I_{24} & I_{25}
 \end{array}
 *
 \begin{array}{ccc}
 w_1 & w_2 & w_3 \\
 w_4 & w_5 & w_6 \\
 w_7 & w_8 & w_9
 \end{array}
 =$$

$$\begin{array}{ccc}
 \begin{array}{c} I_1 w_1 + I_2 w_2 + I_3 w_3 \\ + \\ I_6 w_4 + I_7 w_5 + I_8 w_6 \\ + \\ I_{11} w_7 + I_{12} w_8 + I_{13} w_9 \end{array} &
 \begin{array}{c} I_2 w_1 + I_3 w_2 + I_4 w_3 \\ + \\ I_7 w_4 + I_8 w_5 + I_9 w_6 \\ + \\ I_{12} w_7 + I_{13} w_8 + I_{14} w_9 \end{array} &
 \begin{array}{c} I_3 w_1 + I_4 w_2 + I_5 w_3 \\ + \\ I_8 w_4 + I_9 w_5 + I_{10} w_6 \\ + \\ I_{13} w_7 + I_{14} w_8 + I_{15} w_9 \end{array} \\
 \begin{array}{c} I_6 w_1 + I_7 w_2 + I_8 w_3 \\ + \\ I_{11} w_4 + I_{12} w_5 + I_{13} w_6 \\ + \\ I_{16} w_7 + I_{17} w_8 + I_{18} w_9 \end{array} &
 \begin{array}{c} I_7 w_1 + I_8 w_2 + I_9 w_3 \\ + \\ I_{12} w_4 + I_{13} w_5 + I_{14} w_6 \\ + \\ I_{17} w_7 + I_{18} w_8 + I_{19} w_9 \end{array} &
 \begin{array}{c} I_8 w_1 + I_9 w_2 + I_{10} w_3 \\ + \\ I_{13} w_4 + I_{14} w_5 + I_{15} w_6 \\ + \\ I_{18} w_7 + I_{19} w_8 + I_{20} w_9 \end{array} \\
 \begin{array}{c} I_{11} w_1 + I_{12} w_2 + I_{13} w_3 \\ + \\ I_{16} w_4 + I_{17} w_5 + I_{18} w_6 \\ + \\ I_{21} w_7 + I_{22} w_8 + I_{23} w_9 \end{array} &
 \begin{array}{c} I_{12} w_1 + I_{13} w_2 + I_{14} w_3 \\ + \\ I_{17} w_4 + I_{18} w_5 + I_{19} w_6 \\ + \\ I_{22} w_7 + I_{23} w_8 + I_{24} w_9 \end{array} &
 \begin{array}{c} \\ \\ \\ \end{array}
 \end{array}$$

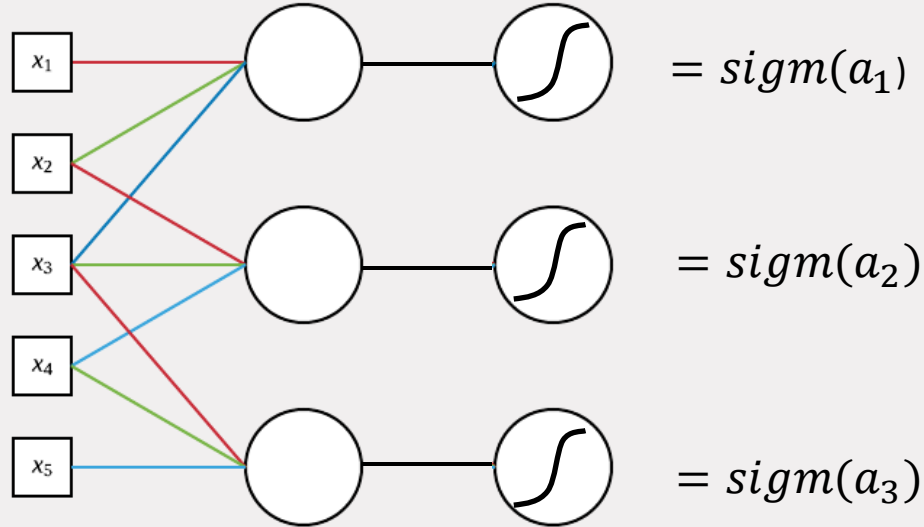
$$\begin{array}{ccccc}
 I_1 & I_2 & I_3 & I_4 & I_5 \\
 I_6 & I_7 & I_8 & I_9 & I_{10} \\
 I_{11} & I_{12} & I_{13} & I_{14} & I_{15} \\
 I_{16} & I_{17} & I_{18} & I_{19} & I_{20} \\
 I_{21} & I_{22} & I_{23} & I_{24} & I_{25}
 \end{array}
 *
 \begin{array}{ccc}
 w_1 & w_2 & w_3 \\
 w_4 & w_5 & w_6 \\
 w_7 & w_8 & w_9
 \end{array}
 + b =$$

$$\begin{array}{ccc}
 \begin{array}{c} I_1 w_1 + I_2 w_2 + I_3 w_3 \\ + \\ I_6 w_4 + I_7 w_5 + I_8 w_6 \\ + \\ I_{11} w_7 + I_{12} w_8 + I_{13} w_9 \end{array} &
 \begin{array}{c} I_2 w_1 + I_3 w_2 + I_4 w_3 \\ + \\ I_7 w_4 + I_8 w_5 + I_9 w_6 \\ + \\ I_{12} w_7 + I_{13} w_8 + I_{14} w_9 \end{array} &
 \begin{array}{c} I_3 w_1 + I_4 w_2 + I_5 w_3 \\ + \\ I_8 w_4 + I_9 w_5 + I_{10} w_6 \\ + \\ I_{13} w_7 + I_{14} w_8 + I_{15} w_9 \end{array} \\
 \begin{array}{c} I_6 w_1 + I_7 w_2 + I_8 w_3 \\ + \\ I_{11} w_4 + I_{12} w_5 + I_{13} w_6 \\ + \\ I_{16} w_7 + I_{17} w_8 + I_{18} w_9 \end{array} &
 \begin{array}{c} I_7 w_1 + I_8 w_2 + I_9 w_3 \\ + \\ I_{12} w_4 + I_{13} w_5 + I_{14} w_6 \\ + \\ I_{17} w_7 + I_{18} w_8 + I_{19} w_9 \end{array} &
 \begin{array}{c} I_8 w_1 + I_9 w_2 + I_{10} w_3 \\ + \\ I_{13} w_4 + I_{14} w_5 + I_{15} w_6 \\ + \\ I_{18} w_7 + I_{19} w_8 + I_{20} w_9 \end{array} \\
 \begin{array}{c} I_{11} w_1 + I_{12} w_2 + I_{13} w_3 \\ + \\ I_{16} w_4 + I_{17} w_5 + I_{18} w_6 \\ + \\ I_{21} w_7 + I_{22} w_8 + I_{23} w_9 \end{array} &
 \begin{array}{c} I_{12} w_1 + I_{13} w_2 + I_{14} w_3 \\ + \\ I_{17} w_4 + I_{18} w_5 + I_{19} w_6 \\ + \\ I_{22} w_7 + I_{23} w_8 + I_{24} w_9 \end{array} &
 \begin{array}{c} I_{13} w_1 + I_{14} w_2 + I_{15} w_3 \\ + \\ I_{18} w_4 + I_{19} w_5 + I_{20} w_6 \\ + \\ I_{23} w_7 + I_{24} w_8 + I_{25} w_9 \end{array}
 \end{array}
 + b$$

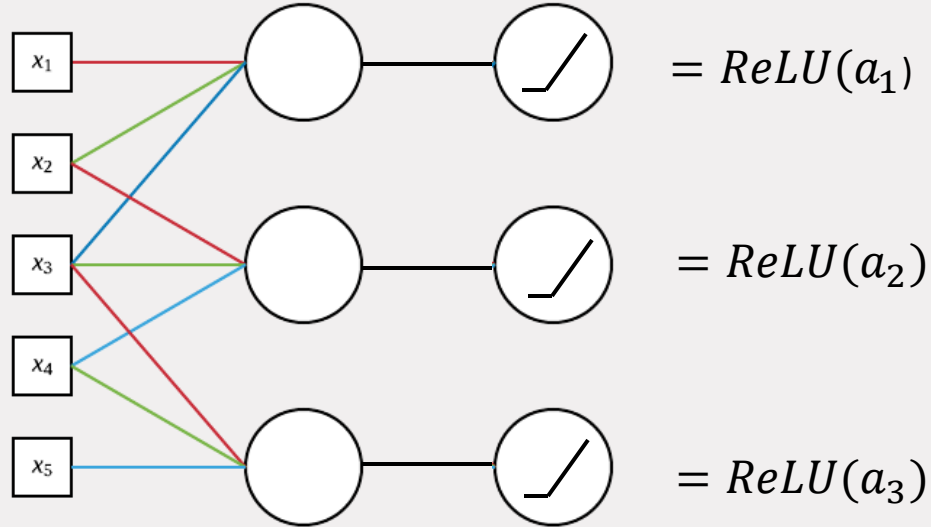
Convolutions + non-linearity



Convolutions + non-linearity (sigmoid)



Convolutions + non-linearity (ReLU)

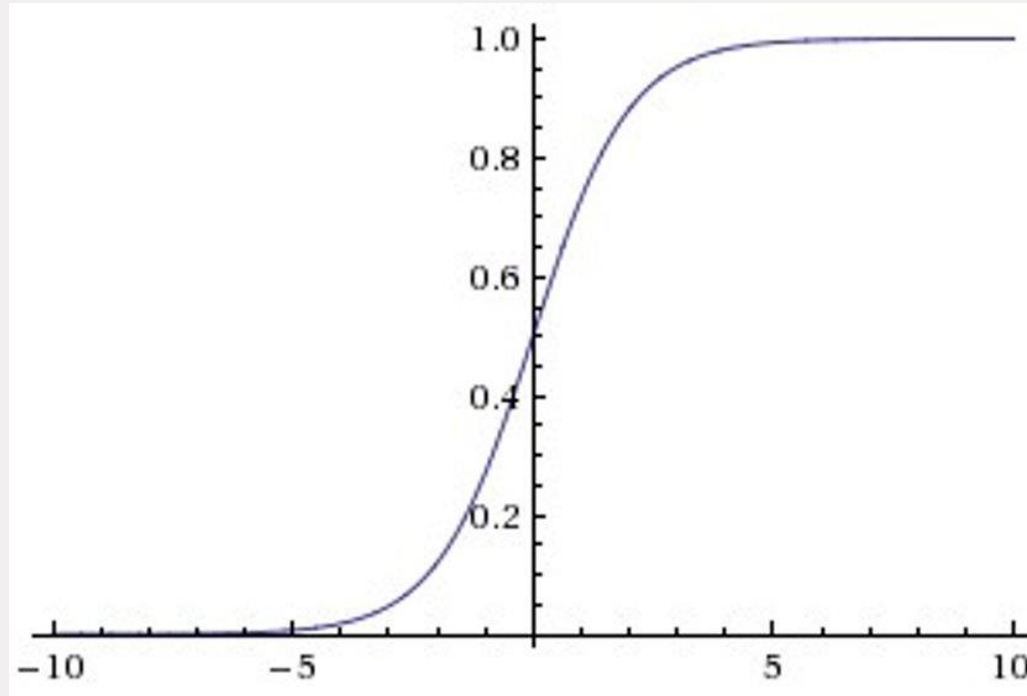


Commonly used activation Functions

- Every activation function (or non-linearity) takes a single number and performs a certain fixed mathematical operation on it.
- There are several activation functions you may encounter in practice.

Credits : <https://cs231n.github.io/>

Sigmoid



Sigmoid

- The sigmoid non-linearity has the mathematical form $\sigma(x)=1/(1+e^{-x})$.
- It takes a real-valued number and “squashes” it into range between 0 and 1.
- In particular, large negative numbers become 0 and large positive numbers become 1.
- The sigmoid function has seen frequent use historically since it has a nice interpretation as the firing rate of a neuron: from not firing at all (0) to fully-saturated firing at an assumed maximum frequency (1).
- In practice, the sigmoid non-linearity has recently fallen out of favor and it is rarely ever used.

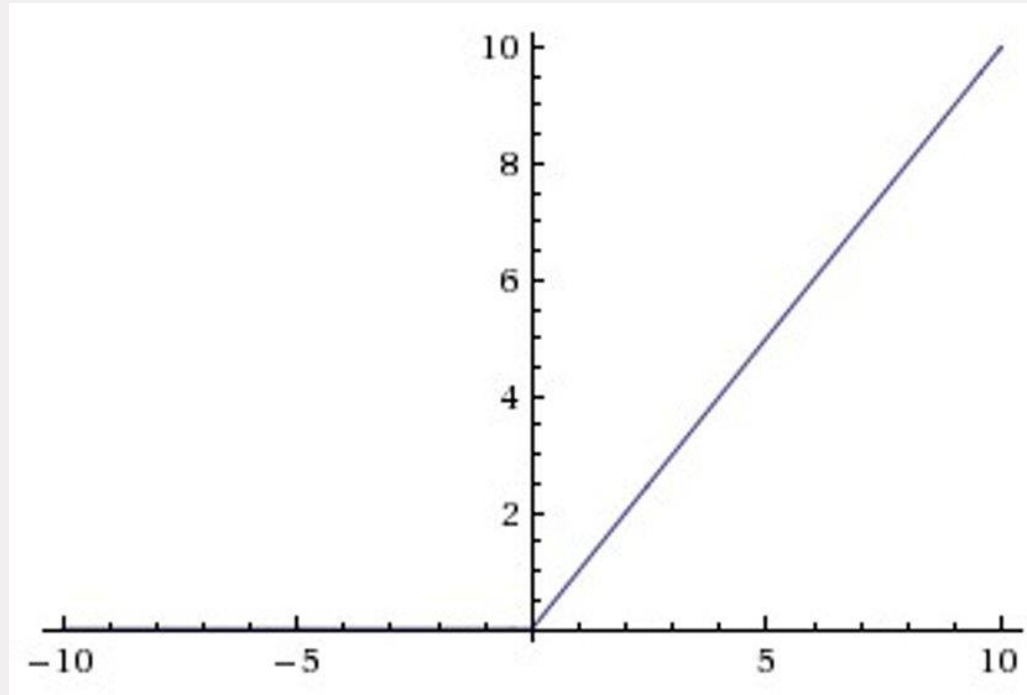
Credits : <https://cs231n.github.io/>

Drawbacks of Sigmoid

- A very undesirable property of the sigmoid neuron is that when the neuron's activation saturates at either tail of 0 or 1, the gradient at these regions is almost zero.
- This (local) gradient will be multiplied to the gradient of this gate's output for the whole objective.
- Therefore, if the local gradient is very small, it will effectively “kill” the gradient and almost no signal will flow through the neuron to its weights and recursively to its data.
- Additionally, one must pay extra caution when initializing the weights of sigmoid neurons to prevent saturation.
- For example, if the initial weights are too large then most neurons would become saturated and the network will barely learn.

Credits : <https://cs231n.github.io/>

ReLU



ReLU

- The Rectified Linear Unit has become very popular in the last few years.
- It computes the function $f(x)=\max(0,x)$.
- In other words, the activation is simply thresholded at zero.

Credits : <https://cs231n.github.io/>

ReLU

- It was found to greatly accelerate (e.g. a factor of 6 in [Krizhevsky et al.](#)) the convergence of stochastic gradient descent compared to the sigmoid/tanh functions.
- It is argued that this is due to its linear, non-saturating form.
- Compared to tanh/sigmoid neurons that involve expensive operations (exponentials, etc.), the ReLU can be implemented by simply thresholding a matrix of activations at zero.

Credits : <https://cs231n.github.io/>

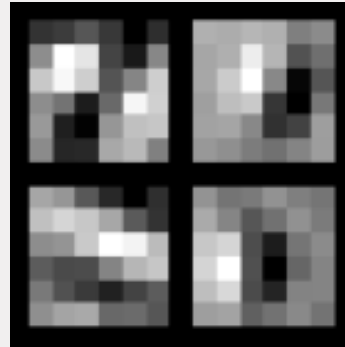
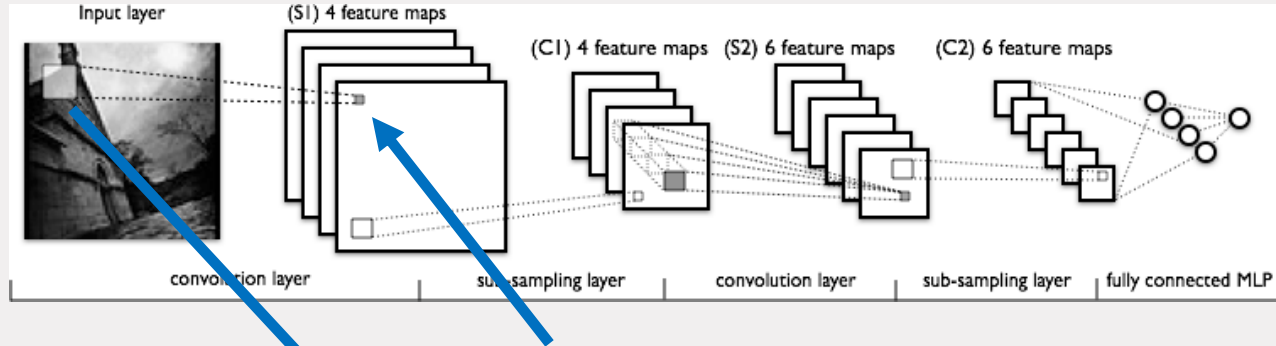
Break!

Quokka (Australia)

<https://imgur.com/r/aww/GqJi0il>



Convolutional kernels



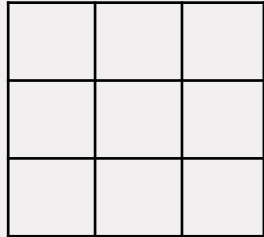
these are the weights
(convolutional kernels)
that produce the four
feature maps

Kernel size

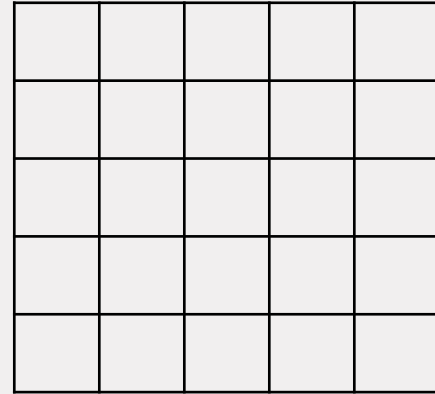
1 x 1



3 x 3

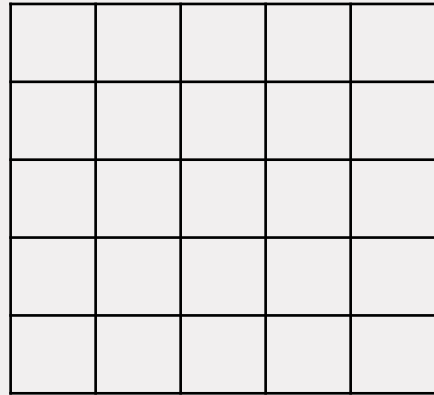


5 x 5



- More weights → more information
- More computations / memory
- Receptive field

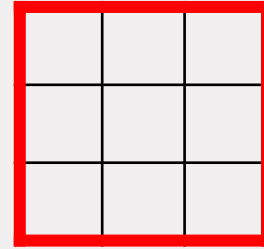
Receptive field



5 x 5

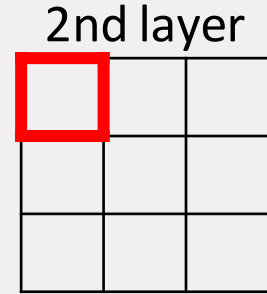
25 weights

Same receptive
field



3 x 3

9 weights



3 x 3

9 weights

= 18 weights

Receptive Field

<https://distill.pub/2019/computing-receptive-fields/>

Computing Receptive Fields of Convolutional Neural Networks

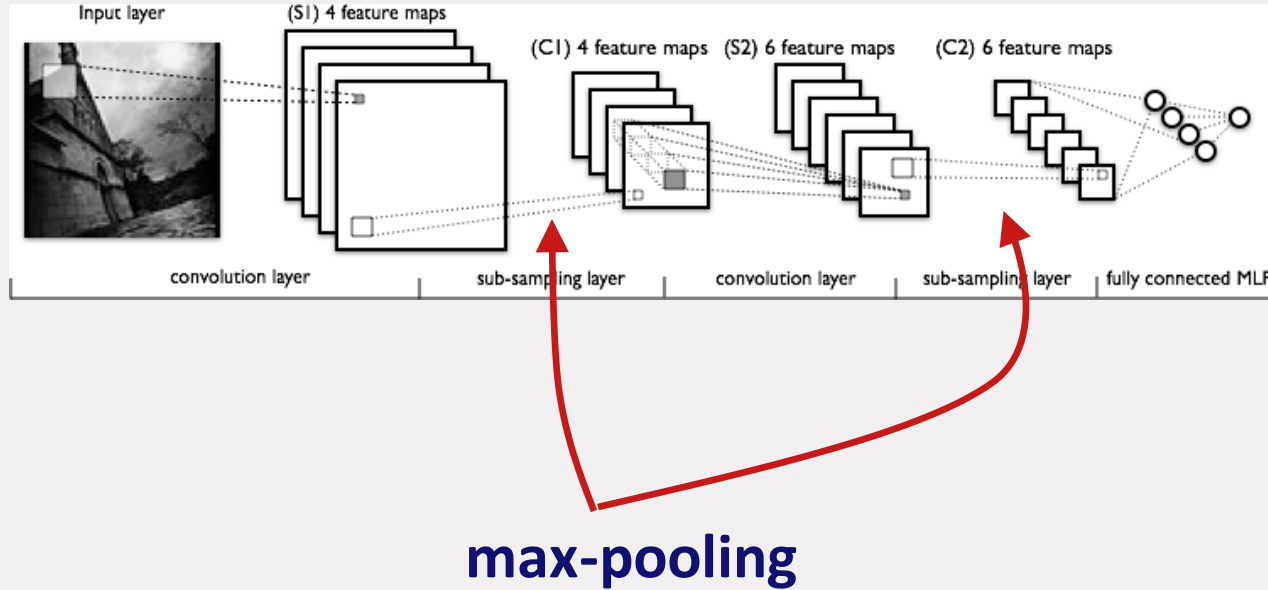
Mathematical derivations and [open-source library](#) to compute receptive fields of convnets, enabling the mapping of extracted features to input signals.

Stride

- We must specify the **stride** with which we slide the filter.
- When the stride is 1 then we move the filters one pixel at a time.
- When the stride is 2 (or uncommonly 3 or more, though this is rare in practice) then the filters jump 2 pixels at a time as we slide them around.
- This will produce smaller output volumes spatially.

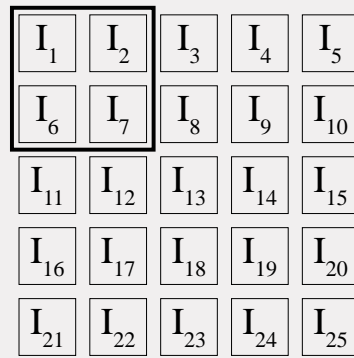
Credits : <https://cs231n.github.io/>

Max-pooling

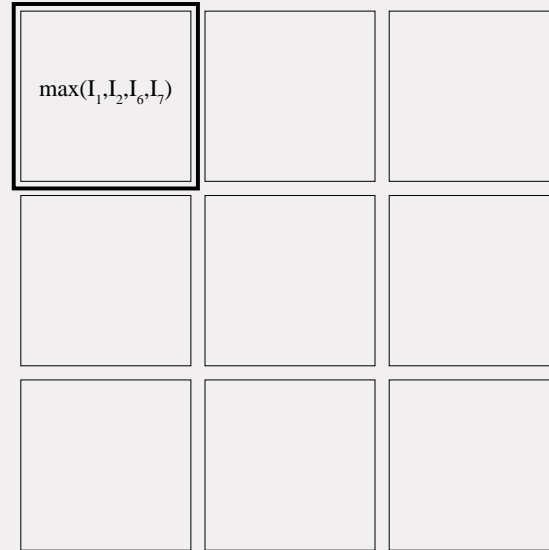


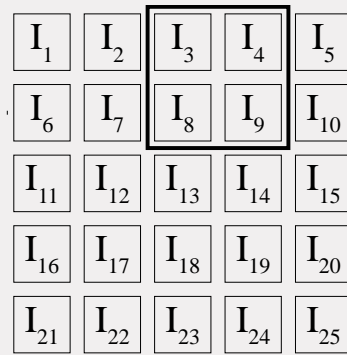
Max-pooling

- Reduce size of feature space
- Maximum of features
- Typical kernel size = 2×2

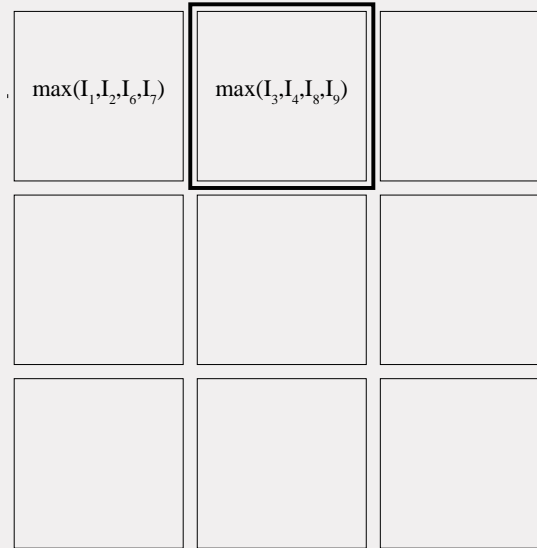


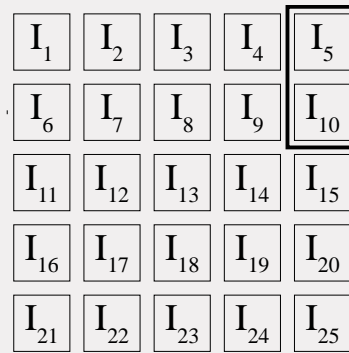
2×2 max pooling



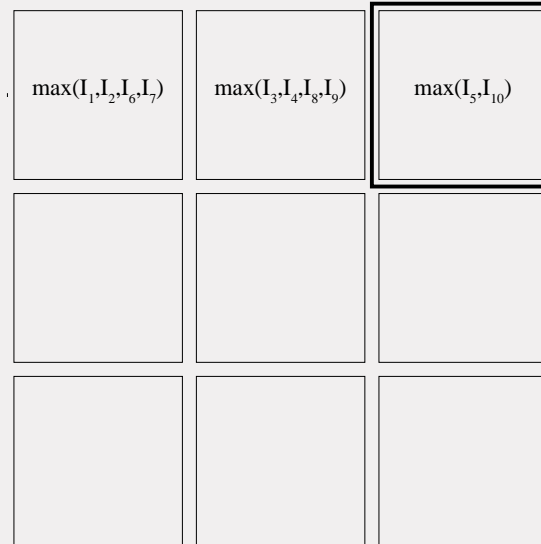


2×2 max pooling

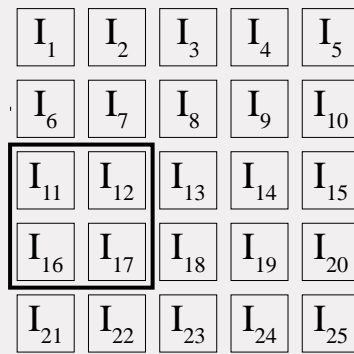




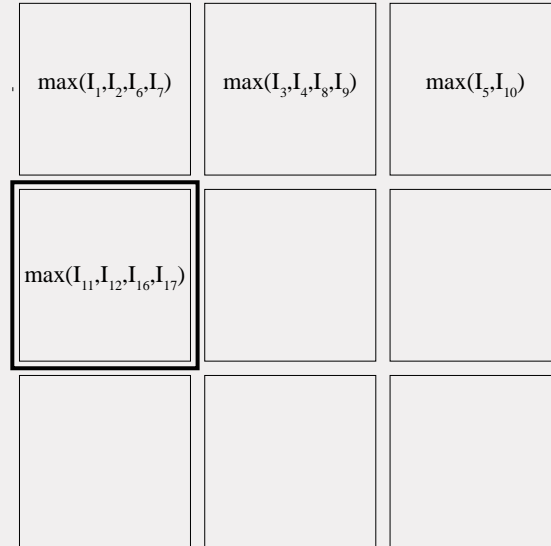
2×2 max pooling

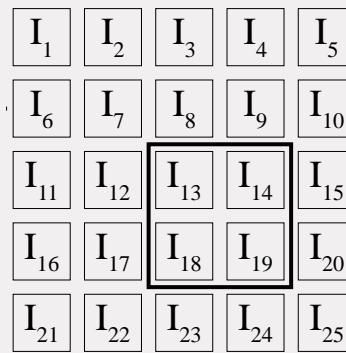


In this example:
no zero padding

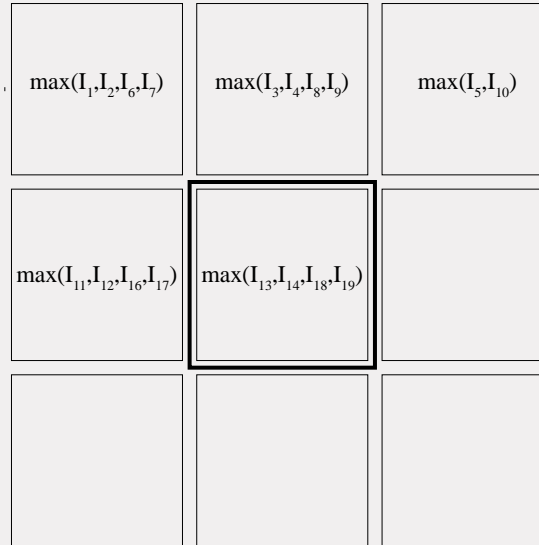


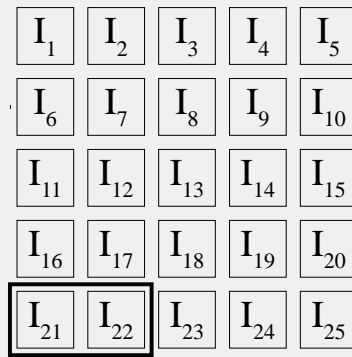
2×2 max pooling



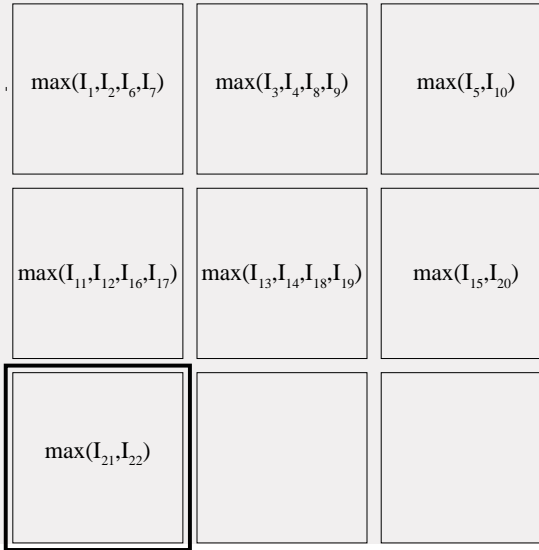


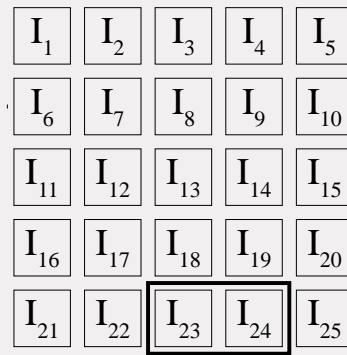
2×2 max pooling



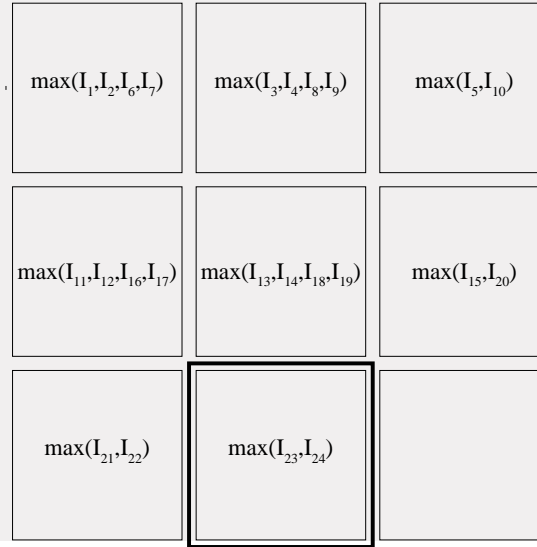


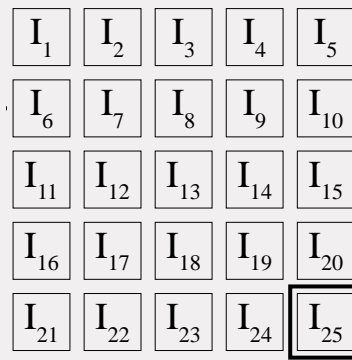
2×2 max pooling



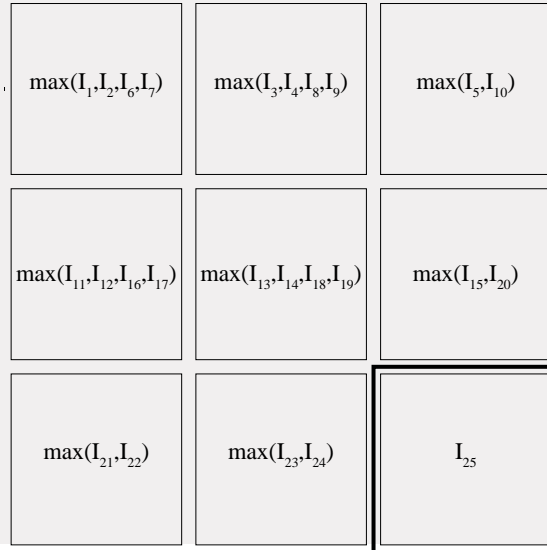


2×2 max pooling

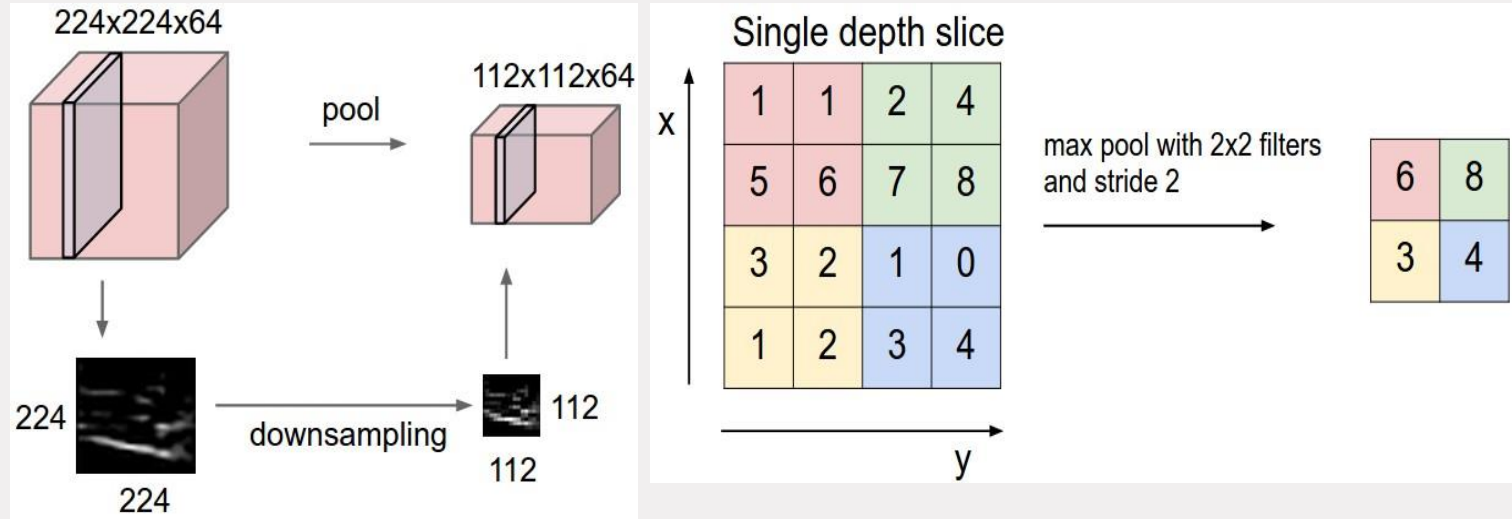




2×2 max pooling



Example:



Pooling layer downsamples the volume spatially, independently in each depth slice of the input volume. **Left:** In this example, the input volume of size $[224 \times 224 \times 64]$ is pooled with filter size 2, stride 2 into output volume of size $[112 \times 112 \times 64]$. Notice that the volume depth is preserved. **Right:** The most common downsampling operation is max, giving rise to **max pooling**, here shown with a stride of 2. That is, each max is taken over 4 numbers (little 2×2 square). <https://cs231n.github.io/>

Benefits of max-pooling

- “Quickly” reduces the number size of the feature maps
- Introduces translation invariance (slightly translated version of the input image will result in the same output)

Alternative: Average Pooling

Demo

<https://poloclub.github.io/cnn-explainer/>

Let try it out!

<http://scs.ryerson.ca/~aharley/vis/conv/flat.html>

An Interactive Node-Link Visualization of Convolutional Neural Networks

Adam W. Harley^(✉)

Department of Computer Science, Ryerson University,
Toronto, ON M5B 2K3, Canada
aharley@scs.ryerson.ca

Training of Convolutional Neural Networks

- Similar to training of 'fully connected' neural networks
- Choose some (random) initial values for network weights
- Optimize networks weights with respect to a loss function that describes difference between network output and label/annotation
- Update networks weights iteratively

Through a process called 'backpropagation'. A good explanation can be found here: <https://www.youtube.com/watch?v=i94OvYb6noo> (5:10 - 28:00, not exam material)

- Keep track of model performance on train & validation set

Summary

- Concept of **convolutions** in a neural network
- Why can we use a **convolutional approach** for (medical) **images**
- Convolutions enable development of **deep** (and large) **neural networks**
- **Max-pooling layer** in a convolutional neural network
- **Kernel size**
- **Receptive field**

